

Sapia.ai — AI Sales Agent Team

Build Walkthrough & Setup Guide

Updated 29 July 2026

Sapia.ai AI Agent Team — Build Complete



Get the repo — two short steps, no memorising git commands.

Step 1: open Claude Code once, anywhere (e.g. Terminal → type `claude` from your home folder or Documents). If you already have a session open, skip to Step 2.

Step 2: prompt it to clone the repo for you:

"Clone <https://github.com/brendanreid-droid/Sales-Agent.git> into my Documents folder, then tell me exactly where it landed."

Claude runs the clone and confirms the folder path. **One thing this can't skip:** a Claude Code session is tied to the folder it started in, so this session won't automatically "become" the Sales Director. Close it, then reopen Claude Code from inside the new folder, double-click `start.command` there, or type `claude` in a

terminal opened to that folder. That second launch is what boots the Sales Director and wires up the 7 specialist agents.

This is a private repo. If the clone fails with a permission error, ask Brendan to add you as a collaborator in GitHub first (repo → Settings → Collaborators).

What Was Built

Everything is now in your **Claude Sale Agent** folder — a complete, production-ready outbound GTM agent system: 7 specialist agents, a tiered target list (plus a continually-growing Lookalike Discoveries table), and a full knowledge vault. Runs on **Claude Code**, that's what lets one orchestrator dispatch all 7 specialists automatically.


Outbound Process Flowchart

End-to-end flow, from new research through to a booked meeting. Blue boxes are agent steps that run on a prompt, pink boxes are the two human checkpoints, nothing crosses either without you.

□ Agent step

□ Human checkpoint

□ Exit / no further action



The dashed line from "Rep reviews in Slack" back to Copywriter is the EDIT path, revise and re-post to the digest.

Files Delivered

8 Agent System Prompts (Sales Director + 7 Specialists) → agent_prompts/

Each file is ready to paste directly into a Claude Project's Custom Instructions field:

File	Claude Project Name
00_sales_director.md	SAPIA-SD — Sales Director
01_research_analyst.md	SAPIA-RA — Research Analyst
02_icp_scorer.md	SAPIA-ICP — ICP Scorer
03_copywriter.md	SAPIA-CW — Copywriter
04_prospect_hunter.md	SAPIA-PH — Prospect Hunter
05_reply_handler.md	SAPIA-RH — Reply Handler
06_company_researcher.md	SAPIA-CR — Company Researcher
07_gtm_action_thinker.md	SAPIA-GAT — GTM Action Thinker

Knowledge Vault → Sapia-Sales-Vault/

Folder	Contents
00-ICP/	ICP Master Profile, Scoring Rubric (7-dimension, weighted), Persona Maps (CHRO, Head of TA, COO, Head of People Systems, CEO, Head of People Services), Target Account Analysis (currently one rep's 155-account territory, each rep replaces with their own, see below), Target Account List CSV
01-Product/	Product Overview, Competitive Differentiators, ROI Data Points
02-Case_Studies/	Case Study Template (with real outcomes reference for all 13 case studies), 13 customer case studies
03-Outreach/	3 email sequences, Subject Line Library, Objection Library (7 objections), Email Winners files
04-Brand_Voice/	Brand Guidelines (tone, prohibited phrases, approved language), Voice Profile Template

Folder	Contents
05-Signals/	Job Movers Log, Newsletter Engagers Log, Trigger Events Log
06- Performance/	Email Performance Tracker, Campaign Learnings, A/B Test Log

Config & Guides

File	Purpose
<code>start.command</code>	Double-click to launch Claude Code (Mac). The only "terminal" step most reps ever need, and this reduces it to one click.
<code>.mcp.json.example</code>	MCP server template for services without a pre-built connector (Path B). No secrets, checked into git, shared automatically when the team clones the repo. Not needed for Slack or Lusha, those connect via <code>claude.ai/customize/connectors</code> instead (Path A).
<code>.mcp.json</code>	Your personal copy with real keys pasted in, only for Path B services. Gitignored, never committed, never shared. You create this by duplicating the file above.
<code>claude_desktop_config_template.json</code>	Only relevant if you're using Claude Desktop instead, see "Why not Claude Desktop?" below. Not needed for Claude Code.
<code>README.md</code>	Full team onboarding guide — share this with your peers

What You Need to Fill In

To ensure business-wide consistency, the central vault has already been fully pre-populated with your verified customer case studies, product overview, and approved ROI database.

Individual Account Executives (AEs) only need to configure their personal outreach settings and API access:

[!IMPORTANT] Do these before running your first agent session:

High Priority (Day 1)

1. **Replace the target list with your own** — `00-ICP/Target_Account_Analysis.md` and `Target_Account_List.csv` currently hold one rep's personal 155-account territory, not shared company data. Swap in your own accounts, same file format, before running anything.
2. **Cold_Winners.md** — AEs add 3–5 of their best-performing personal cold emails (remove prospect names).
3. **Voice Profiles** — Each rep fills in their **Voice_Profile_Template.md** and saves it as `[FirstName]_Voice.md` inside `Sender_Voice_Profiles/`.
4. **Connect your MCP servers** — start with the pre-built connectors at `claude.ai/customize/connectors` (Slack, Lusha, and Gmail, no keys, no file editing). Those three are all you need to run a first batch. Fall back to duplicating `.mcp.json.example` into `.mcp.json` only if you're setting up something else later (SalesLoft, HubSpot, Gong). No terminal needed either way. Full step-by-step below in **"Detailed MCP Connection Setup."**

***On scoring:** whatever ICP score sits in your target list spreadsheet is historical only, ignore it from here on. The only score that counts is what the ICP Scorer agent produces against `00-ICP/ICP_Scoring_Rubric.md`. This keeps every account's priority traceable to one consistent rubric instead of mixing old spreadsheet math with the current one, an account with no ICP Scorer pass yet is unscored, not "whatever the old sheet said."*

GTM Agent Team Roster & Roles

The Sapia GTM team comprises a **Sales Director orchestrating 7 specialized agents** in a coordinated workflow. Here is a summary of each agent's role and primary contribution:

1. 🌀 **Sales Director (SAPIA-SD):** The orchestrator. Coordinates the weekly send schedule, runs the Monday signal harvests, audits Copywriter drafts against the **8-Point Copywriting Checklist**, posts the Slack Morning Digest for approval, and benchmarks Friday performance against real-world metrics.

2. 🔍 **Research Analyst (SAPIA-RA):** The scout. Pulls HubSpot newsletter clicks, monitors job transitions, runs Brave Search lookalike scans to discover net-new companies, and flags triggers.
3. 🏢 **ICP Scorer (SAPIA-ICP):** The qualifier. Automatically scores and tiers prospect accounts (Tier 1/2/3/4) based on a custom 7-dimension rubric to filter out poor-fit leads. Re-scores existing accounts from the target list for consistency.
4. 🔍 **Company Researcher (SAPIA-CR):** The profiler. Builds briefs in two formats, Full (all 10 sections: ATS/HRIS, 4-layer pain points, competitor positioning, the works, for a handful of confirmed Tier A accounts) or Lean (5 sections, batched 3-5 companies per dispatch, for the ongoing weekly research cadence). See "Two Types of Company Research" under Wednesday's cadence step for when to use which.
5. 🗣️ **Copywriter (SAPIA-CW):** The writer. Voice-matches to your reps' profiles and Gong call transcripts, drafting tailored email sequences (Touch 1-4) with high-reply PVP CTAs.
6. 🏹 **Prospect Hunter (SAPIA-PH):** The loader. Enriches contacts with verified email and LinkedIn URLs using Lusha (strictly email-only to protect credit limits), checks 90-day and 21-day cooldowns, and creates the Gmail draft once the batch is approved.
7. 💬 **Reply Handler (SAPIA-RH):** The closer. Monitors replies via Gmail thread tracking, triages intent into 12 buckets, scores BANT criteria, and pushes AE Slack briefs for warm handoffs.
8. 🧠 **GTM Action Thinker (SAPIA-GAT):** The strategist. Challenges, brainstorms, and deconstructs campaign ideas or angles, identifying blind spots and setting metrics-backed roadmaps.

Tuesday & Thursday Cadence Schedule

Nothing on this page runs by itself. There is no background job, no scheduler, no cron. Every single line below happens because you, the AE, open Claude Code and type a prompt at that point in the day. The clock times are a suggested rhythm to aim for, not a timer the system keeps on its own, if you don't open Claude Code and prompt it on Monday morning, Monday's planning simply doesn't happen. Each step below now shows the exact prompt to send.

[MON: Plan & Score] → [TUE: Outreach 1 (Slack Gate)] → [WED: Deep Research] → [

Monday: New Research & Planning

Goal: Set the week's accounts and verify targets.

You open Claude Code and send:

"Let's kick off new research. Run any lookalike discovery we need, and get every candidate account through the ICP Scorer, don't use any legacy score from the target list, only your own current score. Then build the week's outreach plan, capped at 20 unique companies per week, 10 per send day, split across Tuesday and Thursday."

This single prompt is what triggers the Sales Director to dispatch Research Analyst (new lookalikes only) and ICP Scorer (qualify every candidate) in sequence, then assemble the Weekly Outreach Plan. Deliberately kept simple and purely about research and outbounding: newsletter engagers, job movers, and Trigger Events log review are all handled separately (see "Twice Monthly" below), not bundled into every run, for now. It won't happen from a shorter prompt like "good morning", be explicit about what you want done.

Tuesday: Outreach Day 1

Goal: Approve and create Gmail drafts for the first email batch. (Sending itself is a separate, final step you do in Gmail, see below.)

You open Claude Code and send:

"Draft today's Tuesday batch. Use the accounts from Monday's plan, run the Copywriter in sessions of 25, and post the Morning Digest to Slack once every email passes the 8-point audit."

Once the digest is in Slack, you review and reply **APPROVE**, **HOLD**, or an edit note directly in the thread, same as always. Then, back in Claude Code:

"Check Slack for my approval on today's digest and create the Gmail drafts for the approved contacts."

That second prompt is what triggers Prospect Hunter to de-dupe, stage, and now also **automatically create a Gmail draft** for each approved contact, your Slack **APPROVE** is what authorises that, no extra confirmation needed. Claude doesn't poll Slack on its own, you're telling it when to go check. **Drafting is not sending**, no agent ever clicks send, you still open Gmail and send each one yourself, that's the final gate. There's no cadence-tool setting to double-check here, a Gmail draft simply cannot send itself.

Wednesday: Deep Profiling (and New Research, if running it twice a week)

Two different things can happen on a Wednesday, don't confuse them:

1. New Research (optional repeat of Monday), if you want the pipeline topped up mid-week: exactly the same prompt as Monday's "New Research & Planning" above, word for word, same lookalike discovery / ICP scoring step. It's not a different activity that happens to fall on a Wednesday, it's literally Monday's step run again.

2. Deep Profiling (the actual Wednesday-specific step, always runs):

Goal: Build the research briefs for the second send day. There are two different brief formats, pick the right one, see below.

Two Types of Company Research — pick the right one

Using the wrong format either wastes budget on an account that doesn't need it, or under-researches one that does. Confirmed by live testing on 2026-07-15:

	Full (Tailored 1:1)	Lean (Continuous / Bulk)
Use for	A handful of confirmed Tier A accounts entering a real send this week, across all 4 touches	The standing 20-company/week research cadence, exploratory or discovery-stage accounts
What's in it	All 10 sections: snapshot, trigger, ATS/tech stack, full org map (buyer + champion + blocker), 4-layer pain points, competitor mapping, leadership stance, the ONE signal, opener angle	5 sections: snapshot, the single best trigger (not three), buyer/champion (title only unless a contact is already supplied), a one-line pain angle, opener angle
Contact & headcount	Company Researcher researches these directly	Never re-derived by Company Researcher. Supplied as "Given Firmographics" from Prospect Hunter's pre-pull, or marked UNKNOWN and left for Prospect Hunter
Dispatch pattern	One company at a time is fine, quality over throughput	Always batch 3-5 companies into one dispatch, never one at a time, that's where most of the token saving comes from
Confirmed cost	~54k tokens/company (solo)	~18k tokens/company (batched 3-5 per dispatch)

	Full (Tailored 1:1)	Lean (Continuous / Bulk)
Session cap	Max 10 briefs/session	Max 20 briefs/session, run as 4-6 batched dispatches, not 20 individual ones

Full-format prompt (a few high-priority accounts, one at a time is fine):

"Build a Full-format brief for Woolworths Group, this is a confirmed Tier A account going into this week's send. Give me the complete 10-section format, snapshot, trigger, tech stack, full org map with buyer, champion and blocker, 4-layer pain points, competitor mapping, leadership stance, the works. Flag anything too thin for genuine personalisation rather than forcing it."

Lean-format prompt (the ongoing weekly cadence, always batched):

"Build Lean-format briefs for these 5 companies in one batch: [Company A], [Company B], [Company C], [Company D], [Company E]. Read the knowledge base once and reuse it across all five, don't re-read it per company. For headcount and a candidate contact, use whatever I've supplied below as Given Firmographics, and don't re-search for anything I've already given you, if I haven't supplied a contact for one of them, state the target title only and mark the name UNKNOWN for Prospect Hunter to resolve. Spend your search budget on finding the single best trigger per company and the pain/opener synthesis only, don't shop around for a better trigger once you've found one that clears the bar."

If you're doing both New Research and Deep Profiling on the same Wednesday, run New Research first, then Deep Profiling (in whichever format fits) on whatever the Weekly Outreach Plan now includes for Thursday.

Thursday: Outreach Day 2

Same two prompts as Tuesday, swap "Tuesday" for "Thursday" in the first one.

Friday: Audit, Learn & Rubric Calibration

Goal: Collect performance metrics and feed learnings back to the agents.

"Pull this week's stats from Gmail and the cadence tracker, sends, opens, replies, bookings, run the Outbound Analyst Benchmark check, and post the diagnostic

verdict to #brendans-gtm-agent. Then update Campaign_Learnings.md with what worked and what didn't."

Setup Sequence

No "create a Project per agent" step here, the 7 specialists are already wired as native Claude Code subagents in `.claude/agents/`. This is meant to move fast, a rep with their vault populated and MCP servers connected can be running a real batch the same week, most of this is a single afternoon of setup plus one validation pass. No fixed week numbers below, just do them in order.

Core Setup (do these first, in order)

1. Replace the target list with your own territory (`00-ICP/Target_Account_Analysis.md` + `Target_Account_List.csv`)
2. Fill in the rest of the knowledge vault (Cold_Winners.md, other priority files above)
3. Connect Slack, Lusha, and Gmail via the pre-built connectors at `claude.ai/customize/connectors` (Path A below). These three are all you need to get a batch moving, no keys, no file editing. See Detailed MCP Connection Setup below.
4. Fill in your Voice Profile
5. Validate: ask the Sales Director to dispatch Research Analyst, then ICP Scorer, on 1-2 accounts you already know well, check the output quality before trusting it on the full list
6. Run your first full batch: 5-10 emails through the full pipeline → human approval → Gmail send
7. Connect Reply Handler / Slack alerts so replies get triaged automatically once you're sending regularly
8. Settle into the Tuesday/Thursday send cadence

Share the repo with the rest of your team any time after Step 3, each rep repeats Steps 1-4 with their own territory, keys, and Voice Profile, this doesn't block anyone else.

Lower Priority (don't block your first send on these)

These two signal sources run on a set cadence, not continuously, get your core send cycle working first and layer these in once you're comfortable.

- **HubSpot newsletter engagement pull — fortnightly or monthly.** Not yet connected (see MCP Integration Directory below), no-ops until authorised. Feeds the Warm Engagement workflow once it is, including how territory de-duplication works.

- **LinkedIn Sales Nav CSV upload — monthly.** Feeds the Champion Mover workflow.

Ongoing (once the core cycle is running)

- Weekly: review `Campaign_Learnings.md`, update prompts and vault docs based on what's actually landing
- Monthly: calibrate the ICP Scoring Rubric against conversion data
- Begin A/B testing using `A_B_Test_Log.md`
- Consider a third send day if quality and reply rates support it

Optional: Automate the Monday/Wednesday Research Kickoff

You don't have to open Claude Code and type the research prompt yourself every Monday and Wednesday, Claude Code can schedule a recurring cloud routine that does it for you. This isn't a settings panel to click through, you set it up by asking Claude to do it in a normal chat message.

What it does: runs the "new research" step (Lookalike Discovery, ICP Scorer on every candidate) automatically at a set time, writes the results straight into your vault, then posts a short summary to Slack so you know it's ready. Purely research and outbounding prep, Trigger Events, job movers, and newsletter engagers are deliberately kept out of this, see "Twice Monthly" above.

What it deliberately does not do: draft any emails, post a Morning Digest, create any Gmail drafts, or send anything. Those all still need you to open Claude Code and give an explicit prompt, per "Nothing on this page runs by itself" above. This routine only feeds the pipeline, it doesn't act on it.

How to set it up

Open a Claude Code session in this repo and send a prompt like:

"Schedule a recurring routine that runs every Monday and Wednesday at 9am Sydney time. It should act as the Sales Director and do the 'new research' step only: run Lookalike Discovery, and get every new or existing candidate through the ICP Scorer, ignoring any legacy score in the target list. Keep this purely to research and outbounding prep, don't touch Trigger Events, job movers, or newsletter engagers, those are handled separately. Do not draft emails, do not post a Morning Digest, do not create any Gmail drafts. When it's done, post a short summary to [your Slack channel, e.g. #brendans-gtm-agent or a DM] with how many new lookalikes were found and how many accounts were scored."

Claude Code will confirm the schedule back to you before it's active. A few things worth knowing:

- **This is per-rep.** Each rep sets up their own routine pointing at their own vault and territory, there's no shared/central schedule.
- **To check, pause, or change it later,** just ask: *"List my scheduled tasks" or "Pause the Monday/Wednesday research routine."*
- **It still respects every guardrail** (territory dedup, legacy-score ban, lookalike write-back), it's running the exact same "new research" step you'd trigger manually, just on a timer instead of a typed prompt.

Key Design Decisions Applied

Based on your answers, these are the specific configurations built in:

Decision	Implementation
SalesLoft: not yet connected	Planned, but not yet authorised org-wide. Every SalesLoft-shaped step in the agent prompts no-ops until it is; Gmail (below) is the real channel today.
No Salesforce	Removed from all agent prompts and config; <code>Outreach_History_Log.csv</code> + Gmail are the de-dupe sources
Gmail access 	Connected via a pre-built <code>claude.ai/customize/connectors</code> connector, no key, no OAuth file to manage. This is the current send and reply-tracking channel: Copywriter drafts become Gmail drafts, the rep sends every one manually, and thread tracking (<code>thread_id</code>) covers reply detection and follow-up touches.
Gong: not yet connected	Planned for voice-matched calibration against the rep's own calls, not yet authorised. Per-rep Voice Profiles (<code>04-Brand_Voice/Sender_Voice_Profiles/</code>) are the calibration source until it is.
Sales Nav:	Prospect Hunter has full CSV processing workflow; Lusha fills the gaps daily monthly CSV
HubSpot: not yet connected	Planned for a fortnightly/monthly newsletter-engagement query, not yet authorised, no-ops until it is. Research Analyst is briefed on the cadence and territory de-dupe logic ready to go the day it connects.

Decision	Implementation
All outreach via Gmail	Copywriter routes all drafts to Sales Director → Morning Digest → Gmail draft, once the rep approves in Slack. SalesLoft takes over this row automatically once it's connected.
Voice profiles, per-rep	Voice Profile Template created; each rep self-populates. Brendan's is live (04-Brand_Voice/Sender_Voice_Profiles/Brendan_Reid.md), add more as the team grows
Lead product: Chat Interview	Product_Overview.md structured around Chat Interview as primary
Secrets management	<code>.mcp.json.example</code> holds only placeholder values and is safe to commit; each rep's real <code>.mcp.json</code> is gitignored and stays local, no secrets ever touch git

Outbound Email Copywriting Standards

Every outbound email drafted by the Copywriter is audited against a strict, research-backed checklist to ensure it reads like a human-to-human note and maximizes reply rates:

1. The 8-Point Copywriting Checklist

- **Em Dashes (— or –) Ban:** Entirely banned. It makes copy read like a marketing brochure, not a personal email.
- **No Rhetorical Hooks:** Cliches like "Are you struggling with X?" are banned. Opens directly on verifiable triggers.
- **No Generic Flattery:** Empty compliments like "Love what you're building" or "Hope this finds you well" are removed.
- **No Pompous Jargon:** Absolute ban on: *leverage, synergies, cutting-edge, robust, seamless, empower, streamline, unlock, next-level, paradigm, holistic, best-in-class, world-class.*
- **No "I" or "We" Openers:** The email body must start with the prospect's context (e.g. `Given the Q4 hiring spike at [Company]...`).
- **No Meeting Ask in Email 1:** Meeting requests are replaced with soft value CTAs.
- **Strict Word Limits:** Email 1 is capped at 135 words; follow-up touches are capped at 150 words.

- **Sentence-Case Subject Lines:** Capped at 6 words, lowercase, creating curiosity rather than pitching (e.g. `frontline turnover at [Company]`).
- **No Cliché CTAs:** Never "no strings attached." Never a bare "Worth a look?" as its own sentence, fold it into the offer instead: *"Happy to send through [X], if you think it's worth a look?"* Check the sender's Voice Profile for their specific preferred close.
- **Sign-off is name only:** No company name or title line under the sender's name in the email body, the rep's own Gmail signature handles that. Adding one reads as templated.

2. Permissionless Value Promise (PVP) CTAs

Generic "book a call" CTA requests are replaced with high-response **Permissionless Value Promises** (Insight, Benchmark, Playbook, or Diagnostic Questions) that offer standalone value in exchange for a simple "yes" reply: * *Insight CTA:* "Put together a quick one-pager showing how other retail TA teams automate first-round screening. Worth a send?" * *Benchmark CTA:* "Retailers like [Competitor] sit at a 12% candidate assessment drop-off rate. Curious how you compare — happy to share the benchmark data?" * *Diagnostic Question:* "Are you currently using manual resume screening or automated tests first? Reason I ask is that's where candidates usually drop off."

MCP Integration Directory

The agent system is designed around 8 Model Context Protocol (MCP) integrations. As of this update, **Gmail, Lusha, Slack, Brave Search, and Filesystem are connected and live; SalesLoft, HubSpot, and Gong are planned but not yet authorised** org-wide, and any step that depends on them no-ops until they are. The table below maps each integration to the specific agents that use it, their purpose, and when they run.

MCP Server	Agent (Project Name)	Purpose	Operational Timing (When it runs)
Gmail <i>(current send channel)</i>	Sales Director (<code>SAPIA-SD</code>) Prospect Hunter (<code>SAPIA-PH</code>) Reply	Creates the draft for each approved touch (<code>create_draft</code> , never <code>send</code>), searches sent mail and thread history for de-duplication, and tracks replies against each contact's <code>thread_id</code> for the cadence.	Tue/Thu: Draft creation once Slack APPROVE lands. Ongoing: Reply monitoring per the cadence tracker. Friday: Send/open/reply data pull.

MCP Server	Agent (Project Name)	Purpose	Operational Timing (When it runs)
	Handler (SAPIA- RH)		
SalesLoft (not yet connected)	Sales Director (SAPIA- SD) Prospect Hunter (SAPIA- PH) Reply Handler (SAPIA- RH)	Planned to check weekly send volumes, validate existing leads/accounts for de-duplication, stage qualified contacts, and monitor prospect replies natively. Not yet authorised org-wide, every step above no-ops until it is, Gmail (above) fills the same role today.	<i>Not yet live.</i>
HubSpot (not yet connected)	Research Analyst (SAPIA- RA)	Planned to query newsletter click/open analytics for warm intent signals, checked against your own <code>Target_Account_List.csv</code> before being flagged so two reps never chase the same engaged contact off one shared list. Not yet authorised, no-ops until it is.	<i>Not yet live</i> (would run fortnightly/monthly).
Lusha	Research Analyst (SAPIA- RA) Prospect Hunter (SAPIA- PH) Company Researcher (SAPIA- CR)	Lookalike company searches, company size validation, and email-only contact lookup (phone lookups are disabled to save credits).	Monday: Lookalike scans. Tue/Thu: Email-only contact enrichment.

MCP Server	Agent (Project Name)	Purpose	Operational Timing (When it runs)
Gong (not yet connected)	Copywriter (SAPIA-CW)	Planned to query call transcripts strictly filtered by the sender's email address to extract your spoken vocabulary, value messaging, and customer objections. Not yet authorised, per-rep Voice Profiles cover this today.	Not yet live.
Slack	Sales Director (SAPIA-SD) Reply Handler (SAPIA-RH)	Posts regional Morning Digests, weekly plans, Lusha credit alerts, and pushes instant positive reply / meeting AE briefs to reps.	Tue/Thu (8 AM): Pushes digest drafts. Real-time: Pushes meeting/reply notifications (within 5 min).
Brave Search	Research Analyst (SAPIA-RA) Company Researcher (SAPIA-CR)	Web searches for target account news, executive interviews, hiring trends, and ATS keywords (e.g., successfactors logins) in public job postings.	Monday: Lookalike signal checks. Tue/Thu: Generates deep 1-page company research briefs.
Filesystem	All Agents	Reads global ICP profiles, rubrics, persona templates, objection libraries, and brand tone rules. Writes signal logs, briefs, drafts, performance updates, and the outreach history cooldown logs.	Continuously: Interfaced during every active step of the pipeline.

Detailed MCP Connection Setup

There are two different ways a service gets connected here, and most reps will only ever need the first one. **Check the pre-built connector path first**, it's what most people actually

use, the manual `.mcp.json` path is a fallback for whatever doesn't have one yet.

Confirmed 2026-07-15, still current: Slack, Lusha, and Gmail all connect as pre-built connectors, no API key, no `.mcp.json` editing. These three are all a rep needs to run a real batch. SalesLoft, HubSpot, and Gong are not yet authorised org-wide at all, connecting them yourself won't make those steps live, see "Optional, Not Needed Yet" below.

Path A — Pre-built connectors (Slack, Lusha, Gmail)

Anthropic hosts and manages these directly. They're connected once, at your claude.ai account level, not inside this repo, and they then work automatically in every Claude Code session and every project you open, no per-repo setup, no keys to copy.

Step 1 — Check what's already connected

Open Claude Code and send: "Run `/mcp` and tell me which connectors are already live."
Slack is sometimes already connected at the workspace level, if it's already showing as live, you can skip straight to Step 3.

Step 2 — Connect whatever's missing

In a browser, go to `claude.ai/customize/connectors` (log in with the same account you use for Claude Code). Find whichever of **Slack, Lusha, and Gmail** isn't already connected, click **Connect**, and follow the on-screen OAuth sign-in, log into your Slack workspace / Lusha account / Google account and approve. No key to generate, nothing to paste anywhere.

Team / Enterprise plans: connecting a pre-built connector at the account level may require admin permissions. If a connector you need isn't showing up, ask your workspace admin to add it rather than assuming it's broken.

Slack needs one more manual step before it can post anything: connecting the app is not the same as having somewhere for it to post. You still need to create your own private channel and invite the Sapia bot to it, see "Setting Up Slack Notifications & Approvals" below, no MCP tool can do this part for you.

Step 3 — Confirm it's live

Open Claude Code (double-click `start.command` , or type `claude` in a terminal opened to this folder) and send:

"Run /mcp and tell me which connectors are live, and which ones are labelled as coming from claude.ai versus set up locally in this project."

Anything connected via Path A shows up labelled **"from claude.ai"**, that's the pre-built, account-level kind. Anything else in that list came from Path B below, set up specifically for this repo.

Path B — Manual setup, optional, not needed for a first batch

This covers SalesLoft, HubSpot, Gong, and Brave Search, none of which are authorised org-wide yet (see MCP Integration Directory above), and none of which are required to run a first batch, or any batch, right now. Skip this whole section unless you're specifically told one of these has been switched on. This is the older, per-repo, API-key style of connecting, edit one file, no terminal commands beyond opening the app.

One thing to get right before you start: *never paste a real API key into `.mcp.json.example` or commit it to git. Real keys only ever go in your own `.mcp.json` (no ".example"), which is gitignored, stays on your machine, and is never shared with the team.*

Step 1 — Create your personal credentials file

In Finder (or your file browser), find `.mcp.json.example` in the repo root. Duplicate it (right-click → Duplicate, or copy/paste), and rename the copy to exactly `.mcp.json` (drop the ".example"). That's it, no command needed. This new file is already listed in `.gitignore` , so it will never be committed or shared, it's yours alone. Only add entries here for services you're actually setting up this way, don't duplicate Slack, Lusha, or Gmail into this file, they're already handled by Path A.

Step 2 — Get each API key

Service	Where to get your key	Notes
SalesLoft	Your SalesLoft Admin must first enable the Agentic package (Settings → Agentic → Enable). This is an org-wide, one-time step, ask your SalesLoft admin to do this before anyone on the team tries to connect. Once	Not authorised org-wide as of this update, this whole row is future setup.

Service	Where to get your key	Notes
	enabled, each rep generates their own key: Settings → API → Create API Key.	Don't spend time on it until someone tells you the org-wide Agentic package is enabled, Gmail is the working channel until then.
HubSpot	Settings → Integrations → Private Apps → Create a private app. Name it something like "Sapia GTM Agent." Under Scopes, grant read access to Marketing Email and Contacts (for newsletter engagement data). Copy the generated Access Token.	Needs a HubSpot admin if you don't have private-app creation rights. Lower priority, see "Lower Priority" in Setup Sequence.
Gong	Gong → Company Settings → API → Create API Key (generates both an Access Key and a Secret, you need both).	Optional. Only used by the Copywriter for voice calibration against your own past calls, calls are filtered strictly by your own email.
Brave Search	brave.com/search/api → sign up → API key on your dashboard.	Optional. Web search still works without it via Claude's built-in search, this just adds Brave as an extra source.
Filesystem	Not an API key, it's the absolute path to your local copy of <code>Sapia-Sales-Vault</code> , e.g. <code>/Users/yourname/Documents/Sales/Claude Sale Agent/Sapia-Sales-Vault</code> .	Find it by right-clicking the folder → Get Info (Mac) or Properties (Windows) and copying the path.

Step 3 — Fill in .mcp.json

Open your new `.mcp.json` in any plain text editor, TextEdit, Notepad, VS Code, whatever's on your machine already. Find each `"..."` and replace it with the real value you just copied. For example, the SalesLoft block goes from:

```
"env": { "SALESLOFT_API_KEY": "..." }
```

to:

```
"env": { "SALESLOFT_API_KEY": "sk_live_your_actual_key_here" }
```

Do the same for HubSpot, Gong (two values), Brave Search, and the Filesystem path at the very bottom. Leave any `"..."` in place for a server you're not using yet, Claude Code just skips it. Save the file when you're done.

Alternative to editing the file by hand: Claude Code can add one of these servers for you from a chat message instead, useful if a service supports OAuth sign-in rather than a static key:

"Add the SalesLoft MCP server, then run /mcp so I can sign in."

That's the `claude mcp add` command running under the hood, followed by an `/mcp` browser sign-in, you never have to type the command yourself.

Step 4 — Tell Claude to load it

Back in Claude Code, send this message:

"I've filled in my MCP credentials in .mcp.json. Please confirm which servers connected and test each one."

Claude will run through Filesystem, SalesLoft, HubSpot, Gong, and Brave Search (whichever you've filled in) and tell you plainly which ones are live and which ones need attention, no separate test prompts required on your end.

Troubleshooting

- **A pre-built connector (Slack, Lusha, Gmail) isn't showing up in Claude Code:** check `claude.ai/customize/connectors` directly, confirm it shows as connected there first.

On Team/Enterprise plans, only an admin can add one, ask them rather than trying to work around it with a manual API key.

- **Prospect Hunter or Research Analyst return nothing from Lusha when you dispatch them directly by name:** this is a known dispatch quirk, not a broken connection, named specialist agents don't always get tool access from their profile alone. Ask the Sales Director to handle it (it already knows the workaround) rather than dispatching that specialist solo.
- **A manually-configured server shows as not connected:** double check you saved `.mcp.json` (not `.mcp.json.example`) and that the value has no stray quote marks or extra spaces. Ask Claude to re-check after you fix it.
- **SalesLoft key rejected (Path B):** confirm your SalesLoft admin has actually enabled the Agentic package first, this is the most common blocker and it's an org setting, not something you can fix on your own key.
- **Filesystem can't find the vault:** double-check the path in `.mcp.json` is the full absolute path (starts with `/` on Mac) to the `Sapia-Sales-Vault` folder itself, not its parent folder, and that there are no typos.
- **A connector keeps re-prompting for auth:** the OAuth grant may have expired or been revoked on the service's side, just go through the sign-in again, it's harmless to repeat.

Why not Claude Desktop?

Claude Desktop is a genuinely easier product if someone truly cannot use a terminal at all, its Connectors panel is point-and-click with zero typing. But it can't run this system as built: there's no equivalent to `.claude/agents/` or `CLAUDE.md`, so there's no orchestrator automatically dispatching the 7 specialists in parallel, that has to be done by hand, pasting between separate Projects one at a time. It's not the same stack running on a different app, it's a materially more manual process. For a team that needs to work consistently, that outweighs the one-click terminal step Claude Code asks for via `start.command`. If you still need the Desktop path for a specific person, `claude_desktop_config_template.json` is in the repo and uses the same API keys from the table above, just pasted into that file instead and placed at `~/Library/Application Support/Claude/claude_desktop_config.json` (Mac) or `%APPDATA%\Claude\claude_desktop_config.json` (Windows), then restart Claude Desktop.

Sharing with Your Peers

The `README.md` is a complete standalone onboarding guide for new team members. Just share the git repo, `.mcp.json.example` is checked in and holds no secrets, cloning the repo

is the whole "install." Each peer then does their own connector setup from the Detailed MCP Connection Setup above: Path A (their own account-level Slack/Lusha/Gmail connectors at `claude.ai/customize/connectors` , no shared secrets involved at all, and no file editing needed at all for a first batch). Path B (their own `.mcp.json` , their own keys) only matters later, once SalesLoft/HubSpot/Gong are actually authorised.

What each peer needs to do independently:

- Connect their own Slack, Lusha, and Gmail via Path A (a few clicks, no file editing)
- Fill in their own Voice Profile in `04-Brand_Voice/Sender_Voice_Profiles/`

What they get from the shared repo:

- All 7 specialist agents, already wired as native Claude Code subagents, no manual Project setup
- The full Knowledge Vault, including the target account list and the continually-updated Lookalike Discoveries table
- The README setup guide and this walkthrough

Best Practice Guide: Ongoing Operations

This guide provides operational workflows and copy-paste prompts to run the agent team smoothly. One thing to be clear on before any of this: **you talk to a single Claude Code session, one conversation, one "Sales Director."** There's no separate Project per agent to paste into anymore (that was the old Claude Desktop model), the Sales Director reads its identity from `CLAUDE.md` automatically and dispatches the right specialist itself when your prompt calls for it. Your job is to tell it what outcome you want, clearly enough that it knows which agents to involve.

1. Core Prompts You'll Actually Use

These are the same trigger prompts from the Tuesday & Thursday Cadence Schedule above, collected in one place for quick reference, plus the two lower-priority signal workflows.

A. Monday Planning

"Let's kick off new research. Run any lookalike discovery we need, and get every candidate account through the ICP Scorer, don't use any legacy score from the target list, only your own current score. Then build the week's outreach plan, capped at 20 unique companies per week, 10 per send day, split across Tuesday and Thursday."

B. Drafting a Send-Day Batch

"Draft today's [Tuesday/Thursday] batch. Use the accounts from Monday's plan, run the Copywriter in sessions of 25, and post the Morning Digest to Slack once every email passes the 8-point audit."

C. Checking Approvals & Creating Gmail Drafts

"Check Slack for my approval on today's digest and create the Gmail drafts for the approved contacts."

This creates a Gmail draft for each approved contact automatically, it does not send anything. You still open Gmail and send each one yourself when you're ready, that's the final manual step.

D. Twice Monthly: Signal Review, Job Movers & Newsletter Engagers (Lower Priority)

Deliberately separate from the Monday/Wednesday research routine, run this on its own roughly every 1-2 weeks:

"Run a HubSpot newsletter engagement pull. Only flag contacts whose company is in my own Target_Account_List.csv, exclude anything outside my territory, and route any ICP-matched engager to the Warm Engagement workflow."

"I've dropped this month's LinkedIn Sales Nav export at [file path]. Process it for Champion Movers, flag anyone who's moved from a Sapia customer or existing prospect to a new company, and queue new ICP-matched contacts for enrichment."

Anything flagged here rolls into the next Monday/Wednesday's Account Prioritisation step automatically, no need to run it as its own send day.

E. Resuming a Stuck Session

"We got interrupted. Read the latest file status in Sapia-Sales-Vault/ to work out where we left off, give me a brief status summary, and tell me what the next step is before you run any tools."

General rule for all of these: be explicit about the outcome and any constraint that matters (a cap, a cadence, a rule to enforce), a vague prompt like "let's do Monday stuff" won't reliably trigger the full workflow, the Sales Director isn't guessing your intent, it's following what you actually asked for.

2. Setting Up Slack Notifications & Approvals

Reps do not need to check Claude Code constantly. The Sales Director and Reply Handler can push digests and alerts directly to your personal Slack DM or a dedicated channel, and read your response to continue.

Step 1: Authorize Slack

Covered in "Detailed MCP Connection Setup → Path A" above, Slack connects as a pre-built connector at `claude.ai/customize/connectors`, a one-time OAuth sign-in at the account level. It's sometimes already connected at the workspace level, check with `/mcp` first before assuming you need to click Connect yourself. Either way, nothing to configure inside this repo.

Step 2: Create your private channel and invite the Sapia bot

Connecting the app is not the same as having somewhere for it to post, this step is what actually makes posting work, and no MCP tool can do it for you:

- In Slack, create a new private channel named `#[yourfirstname]-gtm-agent`, all lowercase (e.g. `#sarah-gtm-agent`). This naming convention keeps things consistent across every rep's copy of the workspace.
- Invite the Sapia bot/app to that channel (in Slack: open the channel → **Add people** → search for the Sapia app/bot → add it). Digests and alerts will silently fail to post until this is done.

Step 3: Retrieve your Channel or User ID

- **For a channel:** Open Slack, right-click your private channel name (e.g. `#brendans-gtm-agent`, or `#[firstname]-gtm-agent` for your own), select **View channel details**, and copy the **Channel ID** at the bottom (starts with `C`, e.g., `C1234567890`).
- **For personal DMs:** Click your own profile picture in Slack, select **Profile**, click the three dots (`...`), select **Copy member ID** (starts with `U`, e.g., `U1234567890`).

Step 4: Tell the Sales Director where to post

No file editing needed for this one, just tell Claude in chat, e.g. *"Post the Morning Digest to channel C1234567890 and DM me at U1234567890 as well."* It'll use that going forward for the session.

Step 5: The Slack-First Approval Loop

Once you kick off a session in Claude Code:

1. The Sales Director posts the **Morning Digest** (the drafted emails) to your Slack channel.
2. Review the drafts on your phone or Slack client.
3. **Approve directly in Slack:**
 - Reply `APPROVE` to send the whole batch.
 - Reply `HOLD` to pause.
 - Reply `3 EDIT: Change Woolworths hook to focus on retail growth` to revise individual emails.
4. Back in Claude Code, simply type:

"Check Slack approvals and run."

The agent will read the Slack channel history via the Slack MCP, see your approval or edits, update the drafts, and instruct the Prospect Hunter to create the Gmail drafts.

3. Full Guardrail Reference

Every automated check currently built into the system, grouped by what it protects. This is the complete list, not a sample, if it's not here, the agents aren't enforcing it yet.

Human-in-the-Loop & Send Control

Guardrail	Owner	Rule
No autonomous sends, ever	Sales Director / all agents	No agent ever calls a "send" action, full stop. The rep always sends, manually, inside Gmail. The Gmail MCP is only ever used for <code>create_draft</code> and reply-checking, never <code>send</code> .
Auto-draft after Slack approval	Prospect Hunter	Once the rep replies <code>APPROVE</code> in Slack for a batch, approved contacts are staged and a Gmail draft is created for each one automatically, no second confirmation needed. Drafting ≠ sending. No SalesLoft-style Manual/Automatic setting to confirm here, a Gmail draft is inert until the rep opens it and sends it themselves.
No auto-reply	Reply Handler	Drafts every response, never sends. The human rep always sends.

Guardrail	Owner	Rule
Unsubscribe / Negative = immediate stop	Reply Handler	Cadence stopped instantly, contact marked, no reply is sent . No second chances, this one is irreversible by design.

Volume & Batch Caps

Guardrail	Owner	Rule
Weekly batch cap	Sales Director	Max 100 contacts per rep per send day. Overflow queues for the next session.
Unique company limit	Sales Director / Prospect Hunter	Strictly 1 contact per company per batch, the 100 contacts must come from 100 different companies.
Copywriter session size	Copywriter	Runs in sessions of 25 emails, not one session of 100, to manage token limits.
Research session size	Research Analyst	Max 50 accounts per discovery/monitoring session.
Brief session size	Company Researcher	Full format: max 10/session. Lean format: max 20/session, run as 4-6 batched dispatches of 3-5 companies each, never one company per dispatch, see "Two Types of Company Research."
Enrichment session size	Prospect Hunter	Max 100 contacts per session, processed in subsets of 25.

Anti-Spam / Contact & Company Cooldowns

Guardrail	Owner	Rule
90-day contact cooldown	Prospect Hunter	Never re-stage a contact emailed in the last 90 days, regardless of whether they replied. Checked against Gmail sent history and <code>Outreach_History_Log.csv</code> .
21-day company cooldown (new)	Prospect Hunter	Never stage a <i>different</i> contact at a company touched in the last 21 days, even if the earlier contact never replied. Measured from the most recent touch, not the first, so a contact's full 4-touch sequence (~24 days) finishes before anyone else at that company is approached.

Guardrail	Owner	Rule
One active cadence per company	Prospect Hunter	Only one contact per company may be mid-sequence at any time, independent of the day-count above.

Lusha Credit Protection

Guardrail	Owner	Rule
Score-before-enrich	Prospect Hunter	Never run a Lusha lookup on a company that hasn't been scored Tier A or B by the ICP Scorer.
Email-only lookups	Prospect Hunter	Lusha queries fetch email addresses only. Phone number scraping is strictly prohibited, not just discouraged.
No re-enrichment	Prospect Hunter	If a contact already has a verified email logged in <code>Outreach_History_Log.csv</code> , don't spend a credit looking it up again.
No phone fallback	Prospect Hunter	If email verification fails, the contact is flagged HOLD, never enriched with unverified data.
200 credits/rep/week budget	Prospect Hunter / Sales Director	Tracked weekly. Alert posted before any run would push the week's total over budget.

ICP Quality Gate

Guardrail	Owner	Rule
Consistent rubric, no quota-fitting	ICP Scorer	Every company scored the same way regardless of how badly a batch "needs" more accounts.
Boundary ties round down	ICP Scorer	A score that lands exactly on a tier boundary (e.g. 75, 50) defaults to the lower tier.
Never re-prospect a customer	ICP Scorer / Research Analyst	Cross-checked against the confirmed customer list before any account is scored or proposed.
DEI floor	ICP Scorer	Never scores 0 or excludes a company solely for missing DEI evidence, minimum 4/10 on that

Guardrail	Owner	Rule
		dimension regardless.
Unknown ATS = neutral	ICP Scorer	Unconfirmed ATS scores at the midpoint (5/10), not penalised for a data gap.

Research Integrity — No Fabrication

Guardrail	Owner	Rule
Strict UNKNOWN rule	Company Researcher / Research Analyst	Missing data is written as UNKNOWN . Never fabricated, guessed, or extrapolated.
Cite every claim	All research agents	Every fact needs a source URL and a pull date, no exceptions.
No unverified ATS claims	Research Analyst	Never states a company uses a specific ATS unless confirmed via Lusha or a credible public source.
Only verified Sapia capabilities	Copywriter	Only cites features and results that appear in Product_Overview.md / ROI_Data_Points.md . Never invents a feature, a result, or a timeframe that isn't in the source, this includes not attaching an extra qualifier (like an invented timeframe) to an otherwise-real, verified stat.
No personal social links	Research agents	LinkedIn and company profiles only, Facebook/Instagram are banned as sources.

Copywriting Standards (8-Point Audit)

See the full checklist under "Outbound Email Copywriting Standards" above. Every email is checked against it before it reaches the rep.

Lookalike & Target List Hygiene (new)

Guardrail	Owner	Rule
Mandatory write-back	Research Analyst	Every lookalike candidate found gets written into Target_Account_Analysis.md 's Lookalike Discoveries table and Target_Account_List.csv before the session

Guardrail	Owner	Rule
		ends, not just reported. This is what stops the same companies being "discovered" again next week.
Never propose an existing customer	Research Analyst	Checked against the confirmed client list before any lookalike is surfaced.

Reply Handling

Guardrail	Owner	Rule
5-minute escalation	Reply Handler	Positive replies (INTERESTED / INTERESTED-LATER) alerted to the rep within 5 minutes during business hours.
Quote, don't paraphrase	Reply Handler	Slack alerts always include the prospect's literal words.
No follow-up to "not interested"	Reply Handler	A clear decline gets no automated response.
Ambiguous replies escalate, don't act	Reply Handler	When classification is unclear, it's flagged to the rep with reasoning rather than the agent guessing and acting unilaterally.

Data & Credential Protection

Guardrail	Owner	Rule
Salesforce excluded	System-wide	Deliberately left out of every agent prompt and config until your IT team has reviewed and approved it, prevents unreviewed customer-data exposure.
Gong filtered by sender email	Copywriter	Voice calibration queries are strictly filtered to the requesting rep's own calls, never another rep's transcripts.
Secrets never committed (new)	System-wide	<code>.mcp.json.example</code> holds only placeholder values and is safe to commit. Real keys live only in each rep's own <code>.mcp.json</code> , which is gitignored and never leaves their machine.

Brand & Language Standards

Guardrail	Owner	Rule
Australian English only	System-wide	-ise not -ize, -our not -or, per <code>.agents/AGENTS.md</code> .
No "HEXACO"	Copywriter	Say "Sapia's competency framework" in any prospect-facing copy.
No "first mile"	Copywriter	Plain "screening and shortlisting" language instead, American buzzword avoided for an ANZ/APAC audience.
Zero em/en dashes	Copywriter, all agents	Commas, colons, periods, or parentheses instead, applies to every document in the workspace, not just outbound copy.

4. Daily Hygiene Checklist

- **Check Slack first:** Look at `#brendans-gtm-agent` (or your own rep's equivalent private channel) at 9:00 AM local time on Tuesdays and Thursdays.
- **Keep cadences clean:** If a contact replies positively, double-check that the Reply Handler successfully paused their entry in the Gmail cadence tracker.
- **Empty the spam traps:** Review bounces or opt-outs flagged by the Reply Handler once a week to ensure your domain health remains high.
- **Log the wins:** When an email books a meeting, copy the body text and paste it into `Sapia-Sales-Vault/03-Outreach/Email_Winners/Cold_Winners.md` so the Copywriter continually learns from what succeeds.
- **Restart weekly:** If Claude Code feels slow or an MCP connector seems stuck, close the session and relaunch (double-click `start.command` again) to clear connection state.
- **Lusha Budget Lock:** Respect the 200-credit weekly cap. If the Prospect Hunter triggers a warning alert, verify the target list quality before authorizing the override.