

Sapia.ai — AI Sales Agent Team

Implementation Plan & Architecture Guide

Sapia.ai · AI Sales Agent Team · July 2026

Sapia.ai — AI Sales Agent Team: Implementation Plan

Autonomous Outbound GTM Lead Generation Platform on Claude Desktop


Overview

This plan describes a **multi-agent orchestration system** running inside **Claude Desktop** (Team/Enterprise plan) that automates the research, qualification, personalisation, and outreach lifecycle for Sapia.ai's outbound GTM motion. All agents share a common **Foundational Knowledge Vault** and communicate through a **Sales Director orchestrator** that governs cadence, approval gates, and output quality.

The system is designed to be:

- **Human-in-the-loop** at the critical "send" moment — no emails go out without review
- **Cadence-controlled** — 2–3 outreach batches per week, not all-day noise
- **Shareable** — built as Claude Projects so all peers inherit the exact same setup
- **Iterative** — the agents feed learnings back into the knowledge vault over time

1. Platform Architecture



2. Agent Roster — System Prompts & Responsibilities

■ Agent 0: Sales Director (Orchestrator)

Claude Project Name: SAPIA-SD — Sales Director

Role: The conductor. Plans the weekly outreach plan, routes work to specialist agents, sets priorities from signal inputs (newsletter engagement, job moves, new lookalikes), reviews all outputs against the 8-Point Copywriting Audit before surfacing a morning digest to the human SDR for final approval.

Key Behaviours:

- Runs a **Monday planning session** to assign the week's target accounts and outreach slots
- Enforces the **2–3 send-days-per-week** cadence by queuing and scheduling batches
- Audits Copywriter drafts against the **8-Point Copywriting Checklist** (jargon, dashes, openers, CTAs)
- Produces a **"Morning Digest"** each send morning: a Slack-ready summary of all drafted emails ready for human approval
- Flags anomalies: replies, OOOs, negative signals
- Feeds approved/rejected email outcomes back to the Copywriter for learning

■ Agent 1: Company & Stakeholder Research Analyst

Claude Project Name: SAPIA-RA — Research Analyst

MCP Connections: HubSpot, Lusha, LinkedIn Sales Nav*, Brave Search

Role: Deep research on target accounts and stakeholders. Also performs **lookalike discovery** — identifying companies and individuals that match the ICP but are not yet in the database.

Key Behaviours:

- For each target account: industry, size, recent news, tech stack, hiring activity, DEI commitments, ATS/HRIS in use, leadership team
- Identifies key stakeholders: CHRO, Head of Talent Acquisition, CPO, HR Director, COO
- Runs **lookalike scans** using Lusha MCP to find companies matching ICP firmographics (industry, headcount, ATS stack, hiring volume signals)
- Tracks **LinkedIn job moves**: stakeholders from existing customers who have moved to new companies (triggers Champion Mover workflow)
- Flags **HubSpot newsletter openers** who match ICP stakeholder profiles (triggers Warm Engagement workflow)
- Outputs a structured research brief per company in Markdown

■ **Agent 2: ICP Scorer**

Claude Project Name: SAPIA-ICP – ICP Scorer

MCP Connections: HubSpot, Lusha

Role: Objectively scores every prospect company and stakeholder against the Sapia.ai ICP rubric. Prevents wasted outreach on poor-fit accounts.

Scoring Rubric (to be refined with your team):

Dimension	Weight	Signal
Company size (500–50k employees)	20%	Lusha / LinkedIn
Industry fit (Retail/Hosp/Telco/Public Sector)	25%	Research brief
Hiring volume (high-volume, continuous)	20%	LinkedIn jobs, news
DEI mandate evidence	15%	News, website, ESG reports
ATS/HRIS in use (SAP SF, Workday, etc.)	10%	Lusha, BuiltWith
Budget signals (funding round, expansion)	10%	Press, Crunchbase

Output: ICP Score (0–100), Tier (A/B/C), and a one-line rationale per company. Only Tier A and B accounts proceed to outreach.

 ■ **Agent 3: Copywriter / Voice-Matched Writer**

Claude Project Name: SAPIA-CW – Copywriter

MCP Connections: Gong, Filesystem

Role: Drafts all outbound emails. Voice-matches to the individual sender. Calibrates using personal voice profiles + Gong transcripts to build a Customer Language Library. Audits drafts against a strict 8-point copywriting checklist and applies Permissionless Value Promise (PVP) CTAs.

Key Behaviours:

- **Voice Calibration Phase:** Builds customer language libraries from best email winners and Gong calls (filtered strictly by sender's email) before drafting.
 - Enforces **8-Point Copywriting Audit** (no em dashes, no rhetorical hooks, no generic flattery, no "I" openers, no jargon, strict length limits).
 - Implements **Permissionless Value Promise (PVP) CTAs** (Insight, Benchmark, Playbook, Diagnostic Question) replacing meeting pitches.
 - Spaces email sequences on a high-converting rhythm: Day 0 (Opener) / Day 3 (Value benchmark) / Day 7 (Diagnostic question) / Day 12 (Breakup / soft resource).
 - Handles three email archetypes:
 - **Cold outreach** (standard ICP prospect)
 - **Warm engagement** (newsletter opener who is an ICP stakeholder)
 - **Champion mover** (ex-customer stakeholder at new company)
 - Subject lines ≤ 6 words (sentence case, lowercase, no clickbait), body ≤ 120 words for Email 1, ≤ 150 words for follow-ups.
-

■ **Agent 4: SDR / Prospect Hunter**

Claude Project Name: SAPIA-PH – Prospect Hunter

MCP Connections: Lusha, LinkedIn Sales Nav*, SalesLoft

Role: The data acquisition and pipeline loading agent. Finds contacts, enriches records, de-dupes against SalesLoft, and stages approved prospects into SalesLoft cadences.

Key Behaviours:

- Takes the Research Analyst's lookalike list and enriches each company with verified email and LinkedIn URL via Lusha MCP (strictly no phone number lookups to protect credits)
 - Checks SalesLoft for existing contact records to avoid re-prospecting accounts already in an active cadence
 - Stages contacts in SalesLoft (does **not** auto-enrol in a cadence — human confirmation required)
 - Tracks LinkedIn Sales Nav job-move alerts for the Champion Mover workflow
 - Outputs a weekly "Prospect Queue" report with regional tagging (APAC / MENA / UK-EU-US)
-

■ **Agent 5: Reply Handler & Meeting Qualifier**

Claude Project Name: SAPIA-RH – Reply Handler

MCP Connections: SalesLoft, Slack, Filesystem

Role: Monitors replies to outbound sequences using SalesLoft's native reply tracking. Classifies intent, drafts a suggested reply, and qualifies meetings. Escalates positives immediately to the human via Slack — the primary approval and notification channel.

Key Behaviours:

- Classifies replies: Positive / Referral / Objection / OOO / Unsubscribe / Not interested
- For **positives**: drafts a reply proposing meeting time, updates SalesLoft opportunity stage, sends immediate Slack DM alert to the rep with the full draft
- For **objections**: drafts a soft objection-handle reply using the vault's objection library, posts to Slack for rep approval
- For **OOOs**: flags for follow-up on the return date in SalesLoft
- For **unsubscribes**: immediately removes from SalesLoft cadence and logs the action
- After a confirmed meeting: drafts a meeting brief for the AE/rep based on the trigger hook and reply context

3. MCP Server Stack

Tool	MCP Server	Status	Notes
HubSpot	HubSpot Official MCP	■ GA	OAuth 2.0 PKCE; each rep connects their own account filtered by territory
SalesLoft	SalesLoft Native MCP	■ GA	Requires "Agentic" package; admin enables in SalesLoft admin panel
Lusha	Lusha Official MCP	■ GA	API key from Lusha API Hub → https://mcp.lusha.com
Gong	Gong MCP	■ Available	Gong API key (Access Key + Secret) from Gong admin; strictly filtered by sender email
Slack	Slack Official MCP	■ GA	https://mcp.slack.com/mcp — primary approval & notification hub
LinkedIn Sales Nav	Monthly CSV export	■ Safe	Manual monthly export; processed by Prospect Hunter agent
Brave Search	Brave Search MCP	■ Available	Open-web research for Research Analyst
Filesystem	Claude Desktop built-in	■ Built-in	Read/write access to the regional Knowledge Vault folder

Note: Salesforce is intentionally excluded from all MCP connections to prevent any risk of customer data exposure. SalesLoft serves as the single contact management and de-duplication source.

Claude Desktop Config (claude_desktop_config.json) — Template

```
{
  "mcpServers": {
    "hubspot": {
      "type": "http",
      "url": "https://mcp.hubspot.com",
      "env": { "HUBSPOT_ACCESS_TOKEN": "..." }
    },
    "salesloft": {
      "command": "npx",
      "args": ["-y", "@salesloft/mcp-server@latest"]
    },
    "lusha": {
      "type": "http",
      "url": "https://mcp.lusha.com",
      "env": { "LUSHA_API_KEY": "..." }
    },
    "gong": {
      "command": "npx",
      "args": ["-y", "gong-mcp@latest"],
      "env": {
        "GONG_ACCESS_KEY": "...",
        "GONG_SECRET": "..."
      }
    },
    "slack": {
      "type": "http",
      "url": "https://mcp.slack.com/mcp"
    },
    "brave-search": {
      "command": "npx",
      "args": ["-y", "@modelcontextprotocol/server-brave-search"],
      "env": { "BRAVE_API_KEY": "..." }
    },
    "filesystem": {
      "command": "npx",
      "args": [
        "-y",
        "@modelcontextprotocol/server-filesystem",
        "/Users/YourName/Sapia-Sales-Vault"
      ]
    }
  }
}
```

Security Note: Never commit API keys to a shared repository. Use environment variables or a secrets manager (e.g., 1Password CLI, macOS Keychain) and reference them by variable name.

3a. HubSpot Regional Territory Configuration

The team operates across three distinct territories. Each rep's agent team queries HubSpot filtered to their region, ensuring newsletter engagement signals are routed to the correct rep and there is no cross-territory duplication.

Rep Location	Territory	HubSpot Filter Criteria	SalesLoft Queue
■■■ Australia	APAC	Company country IN: Australia, NZ, Singapore, Japan, South Korea, India, SE Asia	APAC-Pipeline
■■■ Dubai	MENA	Company country IN: UAE, Saudi Arabia, Qatar, Egypt, South Africa, rest of Middle East & Africa	MENA-Pipeline
■■■ UK	UK + EU + US	Company country IN: United Kingdom, EU member states, United States, Canada	UK-EU-US-Pipeline

How Territory Filtering Works

HubSpot MCP queries — each rep's Research Analyst agent is instructed to include a territory filter in every HubSpot query. Example instruction in the APAC Research Analyst system prompt:

```
When querying HubSpot for newsletter engagers, always filter by:
  contact.company_country IN [Australia, New Zealand, Singapore, Japan,
    South Korea, India, Philippines, Indonesia, Malaysia, Thailand, Vietnam]
Do not return contacts from other regions – they are managed by separate reps.
```

Lusha lookalike scans — each rep's Prospect Hunter filters by geography:

- APAC rep: company HQ in Asia-Pacific region
- MENA rep: company HQ in Middle East, North Africa, sub-Saharan Africa
- UK rep: company HQ in United Kingdom, Europe, North America

LinkedIn Sales Nav CSV — the monthly export should be pre-filtered by territory before being shared with each rep. Alternatively, each rep exports only their region's saved lists from Sales Navigator.

SalesLoft cadences — set up three parallel cadence groups, one per territory, so send-day timing respects local time zones:

- APAC: send at 8–9am AEST (Tue/Thu)
- MENA: send at 8–9am GST (Tue/Thu)
- UK/EU/US: send at 8–9am GMT/BST (Tue/Thu)

Shared vs. Regional Knowledge Vault

The Knowledge Vault is **shared** across all three reps — same brand voice, same ICP rubric, same product overview, same case studies. Regional variations live in one additional file per territory:

```
/Sapia-Sales-Vault/
■■■ 00-ICP/
  ■■■ Territory_APAC.md      # APAC-specific ICP notes, key accounts, target countries
  ■■■ Territory_MENA.md     # MENA-specific context, cultural messaging notes
  ■■■ Territory_UK_EU_US.md # UK/EU/US regulatory context (GDPR note), key accounts
```

Each rep's agent team reads the global ICP profile **plus** their regional territory file.

4. Foundational Knowledge Vault — Folder Structure

This folder lives on each team member's local machine (or a shared drive) and is mounted via the Filesystem MCP server. It is also uploaded into the shared Claude Project for persistent context.

```
/Sapia-Sales-Vault/
■
■■■■ 00-ICP/
■ ■■■■ ICP_Master_Profile.md          # Firmographic & stakeholder ICP definition
■ ■■■■ ICP_Scoring_Rubric.md          # Scoring criteria & weights (Agent 2 uses this)
■ ■■■■ Persona_Maps.md                # CHRO, HoTA, CPO, COO pain points & priorities
■ ■■■■ Target_Account_List.csv        # Current named accounts
■
■■■■ 01-Product/
■ ■■■■ Product_Overview.md            # What Sapia.ai does, how it works
■ ■■■■ Competitive_Differentiators.md # vs. HireVue, Pymetrics, manual screening
■ ■■■■ Integrations_List.md           # SAP SF, Workday, etc.
■ ■■■■ ROI_Data_Points.md             # Stats, time-to-hire savings, DEI outcomes
■
■■■■ 02-Case_Studies/
■ ■■■■ Retail_Case_Study.md
■ ■■■■ Telco_Case_Study.md
■ ■■■■ PublicSector_Case_Study.md
■ ■■■■ [Additional as available]
■
■■■■ 03-Outreach/
■ ■■■■ Email_Winners/
■ ■ ■■■■ Cold_Winners.md              # Best cold emails with open/reply rates
■ ■ ■■■■ Warm_Engagement_Winners.md   # Emails to newsletter openers that converted
■ ■ ■■■■ Champion_Mover_Winners.md    # Emails to job-movers that converted
■ ■■■■ Subject_Line_Library.md        # Tested subject lines with performance notes
■ ■■■■ Sequence_Templates/
■ ■ ■■■■ Standard_4Touch.md           # Cold email → FU1 → FU2 → Breakup
■ ■ ■■■■ Warm_Engagement_2Touch.md    # Newsletter open → FU
■ ■ ■■■■ Champion_Mover_3Touch.md     # Job move → FU → Breakup
■ ■■■■ Objection_Library.md           # Common objections + suggested responses
■
■■■■ 04-Brand_Voice/
■ ■■■■ Sapia_Brand_Guidelines.md       # Tone, style, prohibited phrases
■ ■■■■ Sender_Voice_Profiles/
■ ■■■■ [YourName]_Voice.md            # Personal writing style reference
■ ■■■■ [Colleague]_Voice.md
■
■■■■ 05-Signals/
■ ■■■■ Newsletter_Engagers_Log.md      # HubSpot openers flagged as ICP (updated weekly)
■ ■■■■ Job_Movers_Log.md              # LinkedIn job-move alerts (updated weekly)
■ ■■■■ Trigger_Events_Log.md          # Funding, expansions, leadership hires
■
■■■■ 06-Performance/
■ ■■■■ Campaign_Learnings.md           # Post-send retrospectives
■ ■■■■ Email_Performance_Tracker.md    # Open, reply, meeting rates by segment
■ ■■■■ A_B_Test_Log.md                # Subject line / hook / CTA tests
```

5. Volume & Cost Management

Design principle: Start small, measure everything, scale on evidence. The system is built to enforce these limits automatically — agents will decline requests that exceed them and explain why.

Starting Volume (Phase 1–2: Weeks 1–4)

Metric	Limit	Reason
Contacts per send session	100	Must be from 100 different companies (1 contact per company per batch) to prevent spamming multiple contacts at the same company simultaneously.
Send sessions per week per rep	2 (Tue + Thu)	200 emails/week — enough signal to learn from, not enough to saturate
Total emails/week per rep	200	Starting volume across both send days
Total emails/week (3 reps)	~600	Manageable for reply monitoring and qualification
Lusha enrichment per week per rep	200 contacts max	Enriches the 200 weekly send-batch — strictly limited to email-only (200 credits total).
Company research briefs per session	10 max	Prevents runaway token usage in a single agent session
Prospect Hunter discovery scan	50 accounts per run	Caps token usage; produces more than enough candidates for the batch

Scaling Protocol — Evidence-Gated

Volume only increases when the data justifies it. The Sales Director reviews these metrics every Friday and posts a scaling recommendation to #sapia-performance.

Scale UP (increase batch to 120 or add a 3rd send day) when ALL of the following are true for 2 consecutive weeks:

- Open rate > 35%
- Reply rate > 5%
- Zero spam complaints or deliverability warnings
- Meeting rate > 1% of sends
- Rep confirms they can handle the increased reply volume

Hold (keep at current volume) when:

- Open rate 25–35% and reply rate 3–5% — system is working, let it compound
- Less than 2 weeks of data — insufficient to make a decision

Scale DOWN (reduce to 50 per batch or drop to 1 send day) if ANY of the following appear:

- Open rate < 20% for 2+ consecutive weeks
- Deliverability warnings in SalesLoft
- Reply rate < 2%
- Domain flagged as spam by any major mail provider
- Rep is unable to action positive replies within the same business day

Do not scale above 150 contacts per batch until the system has been running for at least 8 weeks and has demonstrated sustained performance above all Scale UP thresholds.

Lusha Credit Management

Lusha charges per enrichment (email verification, phone, firmographic data). To avoid wasting credits on low-quality contacts:

Total budget pool: 240,000 Lusha credits.

Credit consumption rule: 1 email address = 1 credit. Phone numbers cost an extra 6 credits.

Always scrape emails only. Scrapes of phone numbers are strictly prohibited.

Rule	Detail
Score before enrich	ICP Scorer must tier a company as A or B <i>before</i> the Prospect Hunter spends any Lusha credits on contact enrichment
Emails only (No phones)	Configure the Lusha MCP query to scrape email addresses only . Do not request phone numbers or dialers. If a phone number is scraped, it consumes 6x extra credits.
No bulk enrichment	Prospect Hunter enriches contacts in batches of 25 — never the full discovery list at once
Weekly credit alert	If estimated credits for the week exceed 200 per rep (600 total), Sales Director posts a warning to #sapia-alerts before proceeding
No re-enrichment	Do not re-enrich contacts already in SalesLoft — check first, enrich only net-new
Skip database gaps	If Lusha returns no email match for a contact, do not purchase a phone credit as a fallback — flag for manual research instead

Claude Token Management

Each agent session consumes Claude tokens. Long research briefs, large discovery scans, and multi-email drafting are the main cost drivers. The following session limits are hard limits enforced by the Sales Director:

Session	Hard Limit	Why
Prospect Hunter discovery scan	50 accounts	Beyond this, a single session risks hitting context limits
Company research briefs (per session)	10 companies	Deep briefs are token-heavy; batch and run across multiple sessions if needed
Email drafting per session	25 emails	Run four Copywriter sessions for a 100-contact batch (25 + 25 + 25 + 25) — better quality, no context overflow
ICP scoring per session	50 contacts	Sufficient for a week's pipeline
Reply Handler per session	Process all replies — no limit	Replies are short; this is low token cost

If you're approaching context limits mid-session: stop, save progress to the vault (Filesystem MCP), start a fresh session, and continue. The Sales Director is responsible for managing this pipeline across sessions.

Cost Warning Triggers (Sales Director Responsibility)

The Sales Director agent must refuse and flag the following requests:

■ GUARDRAIL – Batch size exceeded
Request: [X] contacts for this send session.
Limit: 100 contacts per batch.
Action: Queuing the excess [X-100] contacts for the next send session (Thu/next Tue).
Reason: Sending more than 100 contacts in one batch risks domain reputation and exceeds the agreed review window. Scale up requires 2 weeks of data above thresholds.

■■ LUSHA CREDIT WARNING
Estimated enrichments this week: [X] (limit: 200 per rep)
Current spend: [Y] credits used so far this week
Action: Pausing enrichment. Please confirm you want to proceed beyond the weekly budget.

■■ SESSION TOKEN WARNING
This session has processed [X] research briefs / discovery accounts.
Recommendation: Save progress to vault and continue in a new session to avoid context issues.

6. Weekly Operating Cadence

Monday — Planning & Research

Time	Agent	Action
AM	Sales Director	Pulls HubSpot newsletter engagement data (territory-filtered); identifies ICP-matched openers
AM	Research Analyst	Checks LinkedIn CSV for job-move alerts; processes champion mover candidates
AM	Prospect Hunter	Runs Lusha lookalike scan — max 50 accounts per session
Midday	ICP Scorer	Scores all new accounts; tiers to A/B/C — max 50 contacts per session
PM	Sales Director	Plans the week: selects exactly 100 Tier A/B contacts per send day; queues remainder for following week
PM	Sales Director	Posts weekly Prospect Queue to Slack for rep approval (only new accounts, not re-enriched)

Tuesday & Thursday — Outreach Days

Time	Agent	Action
AM (early)	Copywriter	Drafts batch emails — runs as 4 sessions of 25 emails each to manage token load
AM	Sales Director	Assembles Morning Digest — batch capped at 100 emails — posted to Slack digest channel
~9am	Human rep	Reviews digest in Slack, replies APPROVE / HOLD / [number + edit]
Post-approval	Prospect Hunter	Enriches approved contacts via Lusha (max 100 enrichments per session, run in sets of 25); stages in SalesLoft
All day	Reply Handler	Monitors inbox; surfaces positive replies to human via Slack within 5 min

Friday — Review & Scaling Decision

Time	Agent	Action
AM	Sales Director	Compiles weekly performance summary: sends, opens, replies, meetings booked
AM	Sales Director	Runs scaling check against thresholds (see Section 5); posts recommendation to #sapia-performance
AM	Copywriter	Updates Campaign_Learnings.md with what worked and what didn't
PM	ICP Scorer	Adjusts scoring rubric based on which account tiers responded best

6. Trigger-Based Workflows (Beyond Standard Cadence)

■ Workflow A — Newsletter Engagement (HubSpot Signal)

HubSpot detects newsletter open/click by ICP-matched contact

- Research Analyst: pulls stakeholder brief
- ICP Scorer: confirms Tier A/B fit
- Copywriter: drafts warm "noticed your engagement" email referencing newsletter topic
- Sales Director: adds to next send-day batch
- Human approves → SalesLoft sends

■ Workflow B — Champion Mover (LinkedIn/Lusha Signal)

LinkedIn Sales Nav monthly CSV processed by Prospect Hunter

- Research Analyst: researches new company (ICP fit? already in SalesLoft?)
- ICP Scorer: scores new company
- Copywriter: drafts "congrats on the new role" email leveraging shared history
- Sales Director: marks as Priority – Slack alert to rep within the hour
- Rep approves via Slack → SalesLoft sends within 48hrs of detected move

■ Workflow C — Lookalike Discovery (Weekly Enrichment)

Research Analyst: queries Lusha for companies matching ICP firmographics (filtered by territory)

- ICP Scorer: filters to Tier A/B
- Prospect Hunter: enriches contacts, checks SalesLoft for existing records
- Stages contacts in SalesLoft as "Pending Approval"
- Sales Director: posts weekly Prospect Queue to Slack for rep review
- Rep approves via Slack → contacts enrolled in cadence

■ Agent 6: Company Researcher (Deep-Task Specialist)

Claude Project Name: SAPIA-CR – Company Researcher

MCP Connections: Brave Search, Filesystem

Role: Specialist research agent that takes qualified target accounts and builds detailed, objective, and source-cited 1-page briefs. Maps organization structures, 4-layer pain points, and competitor positioning to feed the Copywriter.

Key Behaviours:

- Builds structured briefs in 60 seconds of reading time (SNAPSHOT, FUNDING, NEWS).
- Maps **4-Layer Pain Points** (Surface, Business, Personal, Career) for key stakeholders.
- Identifies **Competitor Positioning** and failed alternatives the prospect likely evaluated.
- Distills the single most compelling intent or trigger signal from the last 90 days.
- Cites source URLs for every single claim (strict no-fabrication rule).

■ **Agent 7: GTM Action Thinker (Strategist & Challenger)**

Claude Project Name: SAPIA-GAT – GTM Action Thinker

MCP Connections: Filesystem

Role: Senior GTM strategist and Devil's Advocate. Evaluates and deconstructs campaign ideas, angles, and strategies to identify blind spots, test hypotheses, and create execution roadmaps.

Key Behaviours:

- Deconstructs GTM ideas into core hypotheses and categories.
- Audits stated/unstated assumptions, rating them (Validated, Reasonable, Unproven, Risky).
- Runs **Devil's Advocate Challenge** (strongest counterargument, blind spots, uncomfortable question).
- Formulates 10x blueprints, adjacent plays, and contrarian angles.
- Builds Week 1-4 execution roadmaps with leading/lagging metrics and explicit **Kill Criteria**.

7. Human-in-the-Loop Design Principles

*The system is deliberately designed so **no email is ever sent automatically**. Every outreach batch requires explicit human approval. This protects your sender reputation and ensures quality control.*

Gate	What happens	Human action required
Prospect Approval	New lookalike accounts staged in SalesLoft (max 200/week)	Rep reviews weekly Prospect Queue posted to Slack; approves or declines
Morning Digest Approval	All drafted emails posted to Slack by 8am send-day (cap: 100)	Rep reads in Slack, replies APPROVE / HOLD / [number + edit note]
Batch Override Request	Rep requests >100 contacts in a session	Sales Director declines, explains guardrail, queues excess for next session
Lusha Credit Warning	Estimated weekly spend exceeds 200 credits per rep	Rep confirms explicit override before Prospect Hunter continues enrichment
Scaling Recommendation	Friday — Sales Director posts scale-up/hold/scale-down recommendation	Rep confirms or overrides the scaling decision
Positive Reply Escalation	Slack DM alert within 5 min with drafted response	Rep reviews draft in Slack, sends or edits
Objection Response	Draft posted to Slack for review	Rep approves or rewrites before sending
Champion Mover Alert	Slack DM flagged ■ PRIORITY — draft ready	Rep must approve within 48hrs to capture the momentum window
Unsubscribe	Auto-actioned (removed from cadence)	No action required — logged automatically

■ Spam & Overlap Protection Controls

To prevent spamming the same contacts or duplicate outreach to the same companies:

- **90-Day Cooldown:** Prospect Hunter checks SalesLoft and outreach_History_Log.csv. Any contact emailed in the last 90 days is automatically skipped.
- **1 Contact Per Company Per Batch:** Outreach batches are restricted to 1 contact per company at a time. Multi-stakeholder targets are staggered across future send sessions.

8. Slack as the Approval & Notification Hub

Slack is the **primary interface** between the agent team and the human reps. Once the weekly run is started in Claude Desktop, the rep manages all approvals and edits directly inside Slack.

How the Slack Approval Loop Works

- 1. Agent Posts:** Sales Director generates the weekly plan or Morning Digest and posts it directly to your designated regional Slack channel.
- 2. Rep Actions in Slack:** You review the content on your phone/desktop and reply in Slack (e.g. typing APPROVE OR HOLD OR 4 EDIT: make hook shorter).

3. Agent Reads & Executes: When you run the Claude session and say "Check slack approvals and run", the Sales Director calls the Slack MCP tool to fetch your reply, modifies the drafts based on your Slack input, and tells the Prospect Hunter to stage them in SalesLoft.

Slack Channel Setup

Create the following channels in your team Slack workspace:

Channel	Purpose	Who posts
#sapia-digest-apac	Morning Digest + weekly plan for AU rep	Sales Director (AU instance)
#sapia-digest-mena	Morning Digest + weekly plan for Dubai rep	Sales Director (MENA instance)
#sapia-digest-uk	Morning Digest + weekly plan for UK rep	Sales Director (UK instance)
#sapia-alerts	Shared channel — champion movers, urgent flags	All Sales Director instances
#sapia-performance	Weekly performance summaries	All Sales Director instances

How the Approval Flow Works via Slack

Morning Digest (Tue & Thu, ~8am local time):

[illegible]

Positive Reply Alert (within 5 min of detection):

■ POSITIVE REPLY – Woolworths | Sarah Chen, CHRO

"This is timely – we're about to kick off a major hiring ramp.
Happy to jump on a call this week."

Draft response:

"Great to hear, Sarah – Thursday 2pm or Friday 10am AEST work for me.
Happy to adjust. Looking forward to it."

Reply SEND to approve, or paste your edit.

Champion Mover Alert (■ Priority):

■ CHAMPION MOVER – Act within 48hrs

James Wong (ex-Qantas TA Director) → now VP People at Singapore Airlines
Singapore Airlines: APAC, 27,000 employees, high hiring volume – Tier A

Draft email ready. Reply APPROVE to queue for today's send.

Weekly Prospect Queue (Monday):

■ Prospect Queue – Week of 14 Jul

New accounts staged for your approval: 8

Tier A (4): [Company A], [Company B], [Company C], [Company D]

Tier B (4): [Company E], [Company F], [Company G], [Company H]

Reply APPROVE ALL, or list any to HOLD (e.g., "HOLD B, C").

Slack MCP Setup

- Add Slack MCP via Claude Desktop → Customize → Connectors → Search "Slack" → Authorize
- Grant the Slack app permission to post in the designated digest channels
- Each rep configures their own Slack connection (posts to their personal digest channel)
- For DM alerts, the agent posts to the rep's Slack DM directly

9. Claude Projects Setup — Team Sharing

Each agent is its own **Claude Project** within the Team workspace. All projects reference the **shared Knowledge Vault**. Since the team operates across three regions, each rep runs their own instance of the agent Projects — pre-configured with their territory filter and connected to their own Slack channel.

Setup Steps for Each Agent Project:

1. In Claude Desktop → **New Project**
2. Name it per the naming convention: SAPIA-[REGION]-[AGENT] (e.g., SAPIA-APAC-SD)
3. Upload the relevant Vault documents + the rep's regional territory file
4. Paste the agent's system prompt (with territory filter added)
5. Connect the relevant MCP servers for that agent
6. Share the Project with your team via **Team workspace** → **Manage Members**

Regional Project Naming Convention

Agent	APAC (AU)	MENA (Dubai)	UK/EU/US (UK)
Sales Director	SAPIA-APAC-SD	SAPIA-MENA-SD	SAPIA-UK-SD
Research Analyst	SAPIA-APAC-RA	SAPIA-MENA-RA	SAPIA-UK-RA
ICP Scorer	SAPIA-APAC-ICP	SAPIA-MENA-ICP	SAPIA-UK-ICP
Copywriter	SAPIA-APAC-CW	SAPIA-MENA-CW	SAPIA-UK-CW
Prospect Hunter	SAPIA-APAC-PH	SAPIA-MENA-PH	SAPIA-UK-PH
Reply Handler	SAPIA-APAC-RH	SAPIA-MENA-RH	SAPIA-UK-RH

System Prompt Template (add this territory block to each agent):

```
## Territory Configuration
This agent operates for the [APAC / MENA / UK-EU-US] territory.
- HubSpot queries: filter contacts/companies by region (see Territory_[REGION].md)
- Lusha prospecting: filter company HQ by [Asia-Pacific / Middle East & Africa / Europe & North America]
- Slack notifications: post to #sapia-digest-[apac/mena/uk] and DM to [Rep Slack handle]
- SalesLoft cadence: use [APAC / MENA / UK-EU-US]-Pipeline queue
- Send time: [8am AEST / 8am GST / 8am GMT]
```

10. Phased Rollout Plan

Phase 1 — Foundation (Weeks 1–2)

- [] Set up Claude Desktop with Team plan (AU rep leads; Dubai and UK follow)
- [] Build the Sapia-Sales-Vault folder structure; populate all shared documents
- [] Add three regional territory files: Territory_APAC.md, Territory_MENA.md, Territory_UK_EU_US.md
- [] Configure MCP servers (start: HubSpot, SalesLoft, Lusha, Slack, Brave Search)
- [] Set up Slack channels: #sapia-digest-apac, #sapia-digest-mena, #sapia-digest-uk, #sapia-alerts, #sapia-performance
- [] Create SAPIA-[REGION]-SD (Sales Director) and SAPIA-[REGION]-RA (Research Analyst) Projects per rep
- [] Run first Research Analyst session (AU rep) — validate territory filtering and output quality

Phase 2 — Agent Build-out (Weeks 3–4)

- [] Create ICP Scorer Projects per region; calibrate scoring rubric against known customer set
- [] Create Copywriter Projects; each rep populates their Voice Profile
- [] Run first full cold email batch through the system — AU rep goes first as pilot
- [] Test Slack Morning Digest format; refine approval workflow

- [] Share all Projects with Dubai and UK reps; they configure their MCP credentials
- [] Test Newsletter Engagement workflow via HubSpot MCP (territory-filtered query)

Phase 3 — Signal Workflows (Weeks 5–6)

- [] Activate Champion Mover workflow (process first LinkedIn Sales Nav CSV export)
- [] Create Reply Handler Projects per region; connect SalesLoft + Slack per rep
- [] Activate Prospect Hunter; test lookalike discovery with territory filters
- [] All three reps running on 2-day send cadence with Slack approval workflow

Phase 4 — Optimise & Scale (Week 7+)

- [] Monthly: cross-territory performance review via #sapia-performance Slack channel
- [] Calibrate ICP Scoring Rubric per region based on conversion data
- [] Begin A/B testing using the A_B_Test_Log.md
- [] Share learnings across territories — what works in APAC may inform MENA/UK messaging
- [] Consider adding a third send day per region if reply rates justify it

11. Confirmed Decisions Log

All design decisions have been confirmed. This section records the rationale for future reference.

Decision	Choice	Rationale
Salesforce	■ Excluded	Customer data exposure risk — SalesLoft used as sole contact system
Email channel	SalesLoft only	Single tracked, sequenced channel — consistent across all reps
Gmail access	■ Excluded	Not needed. SalesLoft handles all sending and reply tracking natively. Google OAuth setup skipped.
Gong access	Connected ■	Included for voice matched calibration only. Prompts filter calls strictly by sender's email to ensure differentiation.
LinkedIn Sales Nav	Monthly CSV export	Safer from ToS perspective; Lusha fills the real-time enrichment gap
HubSpot engagement	MCP query, fortnightly/monthly	Territory-filtered per rep; no webhook required
SalesLoft package	Agentic package confirmed ■	Native MCP enabled by admin
Team territories	3 reps × 3 regions	AU (APAC), Dubai (MENA), UK (UK+EU+US)
Slack	Primary approval hub ■	All Morning Digests, alerts, and approvals via Slack — reps stay in Slack
Lead product	Sapia Chat Interview	Secondary products (Coaching, Analytics) added to vault as they come in scope
Voice Profiles	3 reps — self-populated	Template provided; each rep fills in their own before first send
Secrets management	TBC with IT	1Password CLI or macOS Keychain recommended